

Stat 610 Homework 5

Due Tuesday, October 20, 11:59pm.

You may work in groups of up to 3 for this assignment.

Assignment

In this assignment, you'll practice code profiling and using git.

You'll start with the code you wrote in the testing lab (<https://jfukuyama.github.io/teaching/stat610/assignments/lab3.pdf>) try out some modifications in an attempt to make it faster.

- Make a github repository, and clone a copy to your computer following the instructions on github. Once this is done, if you type `git status` in the terminal, you should get the output

```
On branch master
```

```
No commits yet
```

```
nothing to commit (create/copy files and use "git add" to track)
```

- Add an initial R file called `llr_functions.R`, testing file called `test_llr.R`, and a file called `benchmark_llr.R`.

Once you've created these files, but before you've added them to the staging area or committed them, if you type `git status`, you should see something like

```
On branch master
```

```
No commits yet
```

```
Untracked files:
```

```
(use "git add <file>..." to include in what will be committed)
```

```
benchmark_llr.R
```

```
llr_functions.R
```

```
test_llr.R
```

```
nothing added to commit but untracked files present (use "git add" to track)
```

Once you're at this stage, add the three files to the staging area using `git add benchmark_llr.R`, `git add llr_functions.R`, `git add test_llr.R`.

At this point, if you type `git status`, you should see something like

```
On branch master
```

No commits yet

Changes to be committed:

(use "git rm --cached <file>..." to unstage)

new file: benchmark_llr.R

new file: llr_functions.R

new file: test_llr.R

Finally, commit these files by typing `git commit -m 'initial commit'`. Feel free to make a better commit message.

Once you have done that, typing `git status` should give you output

On branch master

nothing to commit, working tree clean

- Fill out the `llr_functions.R` and `test_llr.R` files like we did in lab, so that you have a working implementation of `llr` and tests of the functions.

Add these files to the staging area using `git add` and then commit them with `git commit -m 'your commit message'`.

- Add code to the `benchmark_llr.R` file that computes how long your `llr` function takes to run and prints it out (hint: use the function `cat` and one of `bench::mark` or `microbenchmark::microbenchmark`).

Add these files to the staging area using `git add` and then commit them with `git commit -m 'your commit message'`.

Question 1: What does your commit history look like now? What branches do you have, and what commits are they pointing to? Where does HEAD point?

- Make a new branch called `speed-test-1` by typing `git branch speed-test-1`.

You can check that you made the branch by typing `git branch`, which should give you output

```
* master
  speed-test-1
```

which indicates that you have two branches, one called `master` and one called `speed-test-1`, and that `master` is *checked out*, i.e., HEAD points to `master`.

- Check out `speed-test-1` by typing `git checkout speed-test-1`.

Question 2: What changed when you checked out `speed-test-1`? Where does HEAD point now? If you make changes and commit them, where will the `master` and `speed-test-1` branches point?

- *Speeding up llr: take 1:* In the code we wrote in lab, we are using standard matrix multiplication to multiply a diagonal matrix by a dense matrix. If D is a diagonal matrix and X is any

matrix, DX results in the i th row of X being multiplied by D_{ii} , and so there are potentially faster ways of computing DX than the standard matrix multiply. We will try some and see if they actually are faster.

Change the line of code

```
Wz = make_weight_matrix(z, x, omega)
```

so that Wz is a vector of weights instead of a matrix with weights on the diagonal.

Then change the line of code

```
f_hat = c(1, z) %%% solve(t(X) %%% Wz %%% X) %%% t(X) %%% Wz %%% y
```

so that you use the `apply` function in place of `Wz %%% X` and `Wz %%% y`.

Question 2: What function did you use? Why is it equivalent to the matrix multiply?

Add your changes to the staging area using `git add llr.R` and then commit them with `git commit -m 'your commit message'`.

- Use your `benchmark_llr.R` script to check how fast the new version of `llr.R` is. Switch between the master and `speed-test-1` branches by using `git checkout master` and `git checkout speed-test-1` and run `Rscript benchmark_llr.R`.

Question 3: Which version of the function is faster?

- *Speeding up llr: take 2.* Now we'll try another way of speeding up `llr`. Create a new branch for the new version called `speed-test-2` and switch to that branch using

```
git checkout speed-test-1
```

```
git branch speed-test-2
```

```
git checkout speed-test-2
```

Change the line of code

```
f_hat = c(1, z) %%% solve(t(X) %%% Wz %%% X) %%% t(X) %%% Wz %%% y
```

so that you use the `sweep` function instead of `Wz %%% X` and use a vectorized function instead of `Wz %%% y`.

Question 4: What function did you use? Why is it equivalent to the matrix multiplication?

Add your changes to the staging area using `git add llr.R` and then commit them with `git commit -m 'your commit message'`.

- Use your `benchmark_llr.R` script to check how fast the new version of `llr.R` is. Switch between the master, `speed-test-1`, and `speed-test-2` branches by using `git checkout <branch-name>` and running `Rscript benchmark_llr.R`.

Question 5: Which version of the function is the fastest? Do you get a substantial speedup?

- **Question 6:** Run

```
git log --graph --branches
```

What output do you get? What does it tell you about the commit history?

Submission parameters

Submit a pdf with the answers to the six bold-faced questions and a link to your github repository.

Your github repository should have branches **master**, **speed-test-1**, and **speed-test-2**. If John clones your repository, when he checks out the different branches he should see the different implementations of the **llr** function. If he runs `Rscript benchmark_llr.R`, he should get timing information about the version of the function in the branch he's checked out.